

Life Sciences MCP Platform

Grounded AI for Drug Discovery

From researcher question to connected knowledge graph in seconds

The Problem: LLMs Hallucinate Biomedical Knowledge

The status quo is broken:

No Authoritative Access

LLMs rely on stale training data or hallucinate drug-target interactions

Identifier Fragmentation

BRCA1 = HGNC:1100 =
P38398 =
ENSG00000012048 =
672

Manual Integration

Days/weeks per database; code breaks when APIs change

Cognitive Overload

Scientists become integration engineers instead of researchers

*"When asked about drug-target interactions, LLMs either hallucinate or rely on stale training data."
Researchers spend 60-80% of their time on data wrangling*

The Solution: Platform Engineering for Biomedical AI

12 Authoritative APIs in Unified MCP Interface

- **Fuzzy-to-Fact Protocol:** Natural language to canonical identifiers
- **Cross-reference graph:** Automatic multi-hop traversal between databases
- **Token budgeting:** `slim=True` reduces 115+ tokens to 20 tokens per entity
- **Recovery hints:** Error envelopes guide agent self-correction
- **Consistent patterns:** Same schema across all 12 servers

TIER 1: MCP NODES (Verified Entities)

12 FastMCP servers with canonical CURIEs + cross-references

TIER 2: CURL EDGES (Relationship Discovery)

Platform skills for bulk traversal without protocol overhead

TIER 3: GRAPHITI (Research Memory)

Persist validated findings as structured episodes

Quantitative Evidence: Structured APIs Outperform Retrieval

Benchmark Results from 2025 Research

Study	Baseline	With APIs	Improvement
BTE-RAG (GPT-4o mini)	51%	75.8%	+24.8 pp
BTE-RAG (GPT-4o)	69.8%	78.6%	+8.8 pp
RAG Systematic Review	1.0x	1.35x	+35% odds

Key insight: Canonical CURIEs, typed entities, and evidence scores are what benchmarks prove works.

Sources: Xu et al. (2025) "BTE-RAG" bioRxiv; Wang et al. (2025) "RAG in Biomedicine" JAMIA

Value Pillar 1: Rapid API Scaffolding

/scaffold-fastmcp Generates:

Client

Async httpx client with rate limiting

Server

FastMCP server with tool stubs

Models

Pydantic models with CURIE validation

Tests

Integration test scaffolding

Platform Statistics

MCP Servers

12

Total Tests

600+

Client LOC

8,200

Server LOC

2,200

Adding Reactome = 1-2 hours (vs. weeks of custom integration)

Value Pillar 2: Deterministic Nodes, Agentic Edges

The Intentional Separation

MCP SERVERS

12 FastMCP endpoints for verified node retrieval with canonical CURIEs

PLATFORM SKILLS

curl-based workflows for agentic edge discovery without protocol overhead

GRAPHITI MEMORY

JSON episodes to persist validated research findings

Why separate?

MCP carries JSON-RPC protocol overhead (~100ms/call). For bulk edge discovery (10-50 calls/workflow), direct curl is more efficient. Agents choose the right tool for the job.

lifesciences-genomics

lifesciences-proteomics

lifesciences-pharmacology

lifesciences-clinical

The Fuzzy-to-Fact Protocol

Solving the Identifier Problem

PHASE 1 (FUZZY)

"BRCA1" -> Ranked candidates
HGNC:1100 (score: 1.0)
HGNC:1101 (score: 0.8)

PHASE 2 (STRICT)

"HGNC:1100" -> Complete record
ensembl: ENSG00000012048
uniprot: ["P38398"]
entrez: "672"

Why This Matters

- 1 **Prevents hallucination** - Forces entity resolution before retrieval
- 2 **Enables self-healing** - Error envelopes include recovery hints
- 3 **Enables graph traversal** - Cross-references link across databases

```
{"error": {"code": "UNRESOLVED_ENTITY",
"recovery_hint": "Use search_genes()
first to resolve to HGNC CURIE"}}
```

Value Pillar 3: Embracing Our Gaps

Honest Positioning: What We're Not (Yet) Our Philosophy

Gene Ontology Keys

Missing go_process, go_function, go_component

Confidence Calibration

No STRING-style benchmarking against gold standards

Phenotype Integration

No HPO support for phenotype-to-gene queries

DrugMechDB Validation

Gold standard benchmark not yet integrated

1. Document prior art first (sections 1-6 of research docs)
2. Then articulate unique contributions (section 7)
3. Alignment with standards is a **strength**, not a limitation

"Embracing alignment with standards like TRAPI, Biolink, and the Fuzzy-to-Fact pattern means the work builds on proven foundations rather than reinventing them."

Industry Standards Alignment

Standing on Proven Foundations

Standard	Compliance	Notes
TRAPI	85%	Intentional deviations for 60% token reduction
Biolink Model	80%	22-key cross-reference registry
W3C CURIE	Full	Bioregistry-validated prefixes
BioThings Explorer	Aligned	Federated API query patterns

Collaborations

NCATS Translator
TRAPI standard alignment

Bioregistry
Canonical CURIE prefixes

ELIXIR
European life sciences data

Cross-reference registry covers 10/15 major Biolink categories with documented gaps

Novel Contributions

What We Add to the Field

Token Budgeting

Recovery Hints

Fuzzy-to-Fact Protocol

Quality Metrics

CQ Classification

Impact

5-15x context efficiency for batch operations

Autonomous multi-step workflows

Teachable and enforceable patterns

Quantifiable graph quality scores

Systematic capability validation

Key insight: Node retrieval (**MCP**) separated from edge discovery (**curl**) for protocol efficiency

Demo: Synthetic Lethality Target Discovery (ARID1A)

"Show me synthetic lethal partners for ARID1A in ovarian cancer"

- 1 hgnc_search_genes("ARID1A") -> HGNC:11110
- 2 get_gene("HGNC:11110") -> cross_references.ensembl_gene
- 3 string_get_interactions(...) -> EZH2 (0.999), SMARCC1 (0.997)
- 4 chembl_search_compounds("EZH2 inhibitor") -> Tazemetostat
- 5 clinicaltrials_search_trials("Tazemetostat ARID1A") -> NCT03348631

TIME TO INSIGHT

Seconds

BEFORE

1 Week

MANUAL API CALLS

0

ID RECONCILIATION

Automatic

Join Us in Building Grounded Biomedical AI

For Researchers

Use the platform for drug discovery workflows. Test the 15 competency questions in our catalog.

For Bioinformaticians

Contribute API wrappers via /scaffold-fastmcp. Help close the Gene Ontology and HPO gaps.

For AI Engineers

Build agents using MCP tools + platform skills. Explore Fuzzy-to-Fact for other domains.

For Standards Bodies

Help us align with TRAPI, Biolink, DrugMechDB. Validate our cross-reference registry.

Canonical identifiers. Structured responses. Grounded AI.

REPOSITORY

[lifesciences-research](#)

DOCUMENTATION

[docs/platform-engineering-rationale.md](#)

SPECIFICATION

[docs/adr/accepted/adr-001-v1.4.md](#)